

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from math import e
from math import log
import scipy.io as sc
from scipy.special import expit, logit
%matplotlib inline
```

```
In [2]: def logistic_fn(x,w):

    return expit(np.dot(x,w))
```

```
In [3]: def update(x,w,y,s_vals):
    hess_inv = np.linalg.inv(np.matmul(np.matmul(x.T,np.diag(s)),x))
    diff = np.matmul(x.T,(y - s))
    return w - np.matmul(hess_inv,diff)
```

Problem 2.4 (do not submit)

```
In [100... x = np.array([[0.2, 3.1, 1],
               [1.0, 3.0, 1],
               [-0.2, 1.2, 1],
               [1.0, 1.1, 1]])
```

```
In [101... y = np.array([1,1,0,0])
```

```
In [108... w0 = [-1, 1, 0]
```

```
In [109... s0=[]
for i in range(x.shape[0]):
    s0+=s(x[i],w0)
```

```
In [119... s0 = logistic_fn(x,w0)
```

```
In [120... s0
```

```
Out[120]: array([0.94784644, 0.88079708, 0.80218389, 0.52497919])
```

```
In [111... w1 = update(x,w0,y,s0)
```

```
In [112... w1
```

```
Out[112]: array([-1.09534796, 0.43773381, 1.67830773])
```

```
In [113... s1 = logistic_fn(x,w1)
```

```
In [118... s1
```

```
Out[118]: array([0.94354568, 0.86945643, 0.91853884, 0.74354326])
```

Q3.2

```
In [6]: from sklearn.model_selection import KFold
from sklearn.metrics import roc_auc_score
from sklearn.model_selection import train_test_split
```

```
In [338... #returns the mean and stdev of the features of the input data as a list of tuples
```

```
def get_fit(data):

    stats = []

    for i in range(data.shape[1]):

        stats+=[(data[:,i].mean(), data[:,i].std())]

    return stats
```

```
In [339... #normalizes the features of a dataset by subtracting the training mean and dividing by
#should be a list of tuples
```

```
def transform(data,stats):

    for i in range(data.shape[1]):

        data[:,i] = (data[:,i]-stats[i][0])/stats[i][1]

    return data
```

```
In [340... #update function for gradient descent with l2 penalty
```

```
def logr_update(data, labels, s_vals, weights, step, penalty):

    return weights+step*(np.matmul(data.T,(labels - s_vals)) - penalty*weights)
```

```
In [341... #append a column of ones for the fictitious dimension
```

```
def add_fic(data):

    if data.ndim ==1:

        data=data.reshape(1,data.shape[0])

    return np.append(data,np.ones((data.shape[0],1)),axis=1)
```

```
In [342... #calculates the cost of the logistic regression cost function with l2 penalty
```

```
def logr_cost(data, labels, s_vals, weights, penalty):

    log_s = [log(x+1e-10) for x in s_vals]

    log_s_comp = [log((1-x)+1e-10) for x in s_vals]

    return -np.dot(labels,log_s)-np.dot((1-labels),log_s_comp)+penalty*np.linalg.norm(weights)
```

```
In [348... # Define the decision function for
```

```
def predict(x, weights):
```

```
#calculate the logistic function with the  
score = logistic_fn(x,weights)
```

```
return (score >= 0.5).astype(int)
```

In [349...

```
#Performs batch gradient descent with a fixed number of iterations  
def bgd(data, labels, step, penalty,iterations):  
  
    #initialize the weight with the zero vector  
    weight = np.zeros(data.shape[1])  
  
    s_vals = logistic_fn(data,weight)  
  
    #initialize the cost with the first weight vector of all zeros  
    cost=[logr_cost(data, labels, s_vals, weight, penalty)]  
  
    for i in range(iterations):  
  
        #creat the w' vector for the update and calculating the cost  
        w_p = np.append(weight[:-1],0)  
  
        weight = logr_update(data,labels,s_vals,w_p, step, penalty)  
  
        cost+=[logr_cost(data, labels,s_vals, w_p, penalty)]  
  
        s_vals = logistic_fn(data,weight)  
  
    return weight,cost
```

In [350...

```
data=sc.loadmat('data.mat')
```

In [351...

```
labels=data['y'].reshape(len(data['y']))
```

In [352...

```
test = data['X_test']
```

In [353...

```
x = data['X']
```

In [354...

```
k_fold = KFold(n_splits=10, random_state = 42, shuffle=True)
```

In [355...

```
steps = [0.1, 0.01, 0.001, 0.0001] # step size values to try  
lambdas = [10, 1,0.1, 0.01] # l2 penlaty values to try
```

In [356...

```
# Iterate through each combination of hyperparameters to optimize the model  
for step in steps:  
    for penalty in lambdas:  
  
        # Initialize list for cross-validation scores  
        cv_costs = []  
        cv_auc_roc = []  
  
        # Iterate through each fold of the data  
        for train_indices, val_indices in k_fold.split(x_train):  
  
            # Split the data into training and validation sets  
            X_train, y_train = x[train_indices], labels[train_indices]
```

```

X_val, y_val = x[val_indices], labels[val_indices]

# normalize the training and validation data by the training data values
fit = get_fit(X_train)
X_train = transform(X_train, fit)
X_train = add_fic(X_train) #add column of 1s for the bias

X_val = transform(X_val, fit)
X_val = add_fic(X_val) #add column of 1s for bias

weights, cost = bgd(X_train, y_train, step=step, penalty=penalty, iterations=iterations)

# Calculate the cost function on the validation set

s_vals = logistic_fn(X_val, weights)
cost_val = logr_cost(X_val, y_val, s_vals, weights, penalty=penalty)

# Make predictions on the validation set
y_pred = predict(X_val, weights)

# Calculate the AUC-ROC score on the validation set
auc_roc = roc_auc_score(y_val, y_pred)

# Append the cross-validation scores to the lists
cv_costs.append(cost_val)
cv_auc_roc.append(auc_roc)

# Compute the mean and standard deviation of the cross-validation scores
cv_mean_cost = np.mean(cv_costs)
cv_std_cost = np.std(cv_costs)
cv_mean_auc_roc = np.mean(cv_auc_roc)
cv_std_auc_roc = np.std(cv_auc_roc)

print(f'step size={step}, lambda={penalty}, mean CV cost={cv_mean_cost:.3f}, s

```

```

step size=0.1, lambda=10, mean CV cost=6231.396, std=211.044, mean CV AUC-ROC=0.702,
std=0.045
step size=0.1, lambda=1, mean CV cost=535.677, std=305.692, mean CV AUC-ROC=0.963, st
d=0.024
step size=0.1, lambda=0.1, mean CV cost=301.604, std=178.119, mean CV AUC-ROC=0.974,
std=0.009
step size=0.1, lambda=0.01, mean CV cost=154.027, std=91.185, mean CV AUC-ROC=0.983,
std=0.007
step size=0.01, lambda=10, mean CV cost=239.888, std=105.991, mean CV AUC-ROC=0.958,
std=0.036
step size=0.01, lambda=1, mean CV cost=38.726, std=12.382, mean CV AUC-ROC=0.988, std
=0.004
step size=0.01, lambda=0.1, mean CV cost=33.457, std=13.728, mean CV AUC-ROC=0.985, s
td=0.007
step size=0.01, lambda=0.01, mean CV cost=34.857, std=9.045, mean CV AUC-ROC=0.983, s
td=0.004
step size=0.001, lambda=10, mean CV cost=100.352, std=11.567, mean CV AUC-ROC=0.978,
std=0.006
step size=0.001, lambda=1, mean CV cost=53.272, std=11.855, mean CV AUC-ROC=0.981, st
d=0.004
step size=0.001, lambda=0.1, mean CV cost=44.340, std=12.029, mean CV AUC-ROC=0.980,
std=0.004
step size=0.001, lambda=0.01, mean CV cost=43.333, std=12.052, mean CV AUC-ROC=0.980,
std=0.004
step size=0.0001, lambda=10, mean CV cost=111.886, std=12.561, mean CV AUC-ROC=0.966,
std=0.004
step size=0.0001, lambda=1, mean CV cost=71.700, std=12.772, mean CV AUC-ROC=0.967, s
td=0.004
step size=0.0001, lambda=0.1, mean CV cost=67.046, std=12.796, mean CV AUC-ROC=0.967,
std=0.004
step size=0.0001, lambda=0.01, mean CV cost=66.573, std=12.799, mean CV AUC-ROC=0.96
7, std=0.004

```

Based on the above metrics there are a couple of competing combinations to check based on a low cost and high AUC:

step size=0.01, lambda=1, and step size=0.01, lambda=0.1

```
In [357... x_train, x_val, y_train, y_val = train_test_split(x, labels, test_size = 0.2, random_st
```

```
In [358... #normalize the data
fit = get_fit(x_train)
x_train = transform(x_train, fit)
x_train = add_fic(x_train) #add column of 1s for the bias

x_val = transform(x_val, fit)
x_val = add_fic(x_val) #add column of 1s for bias

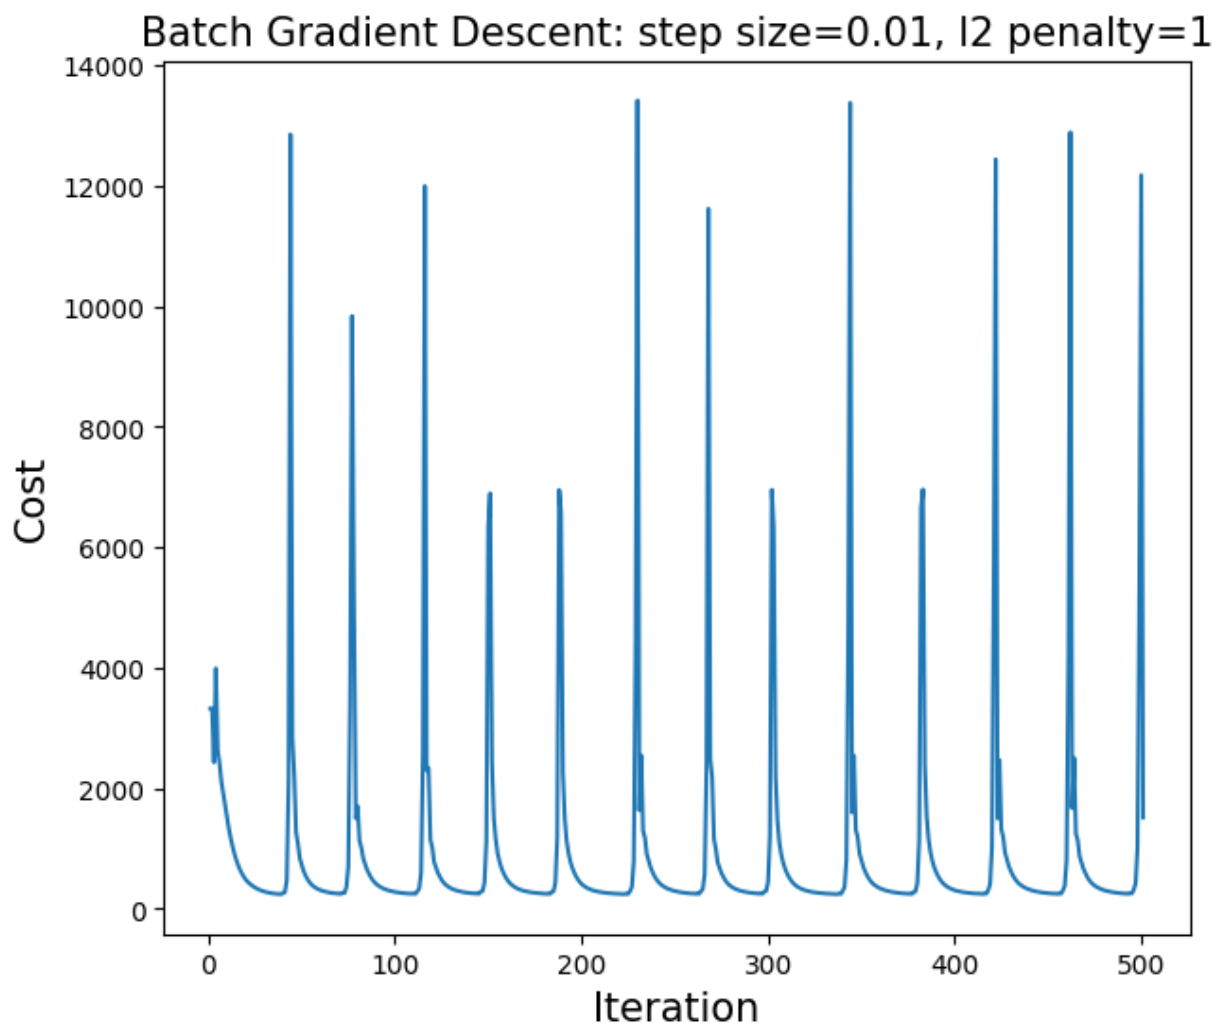
```

```
In [359... w_final, cost = bgd(x_train, y_train, 0.01, 1, 500)
```

```
In [360... plt.figure(figsize=(7,6))
plt.plot(np.linspace(1,501,501), cost)
plt.xlabel('Iteration', fontsize=15)
plt.ylabel('Cost', fontsize=15)
plt.title('Batch Gradient Descent: step size=0.01, l2 penalty=1', fontsize=15)
# plt.savefig('Q3.2 0011.png', dpi=300)

```

Out[360]: Text(0.5, 1.0, 'Batch Gradient Descent: step size=0.01, l2 penalty=1')




This has some strange oscillations. Let's not go with that.

```
In [361...] w_final, cost = bgd(x_train, y_train, 0.01, 0.1, 500)
```

```
In [362...] plt.figure(figsize=(7,6))
plt.plot(np.linspace(1,501,501),cost)
plt.xlabel('Iteration', fontsize=15)
plt.ylabel('Cost', fontsize=15)
plt.title('Batch Gradient Descent: step size=0.01, l2 penalty=0.1', fontsize=15)
# plt.savefig('Q3.2 00101.png',dpi=300)
```

Out[362]: Text(0.5, 1.0, 'Batch Gradient Descent: step size=0.01, l2 penalty=0.1')

Batch Gradient Descent: step size=0.01, l2 penalty=0.1




Still has oscillations, but not as bad. Let's test some other lambdas

```
In [363... w_final, cost = bgd(x_train, y_train, 0.01, 0.01, 500)
```

```
In [364... plt.figure(figsize=(7,6))
plt.plot(np.linspace(1,501,501),cost)
plt.xlabel('Iteration', fontsize=15)
plt.ylabel('Cost', fontsize=15)
plt.title('Batch Gradient Descent: step size=0.01, l2 penalty=0.01', fontsize=15)
# plt.savefig('Q3.2 001001.png',dpi=300)
```

```
Out[364]: Text(0.5, 1.0, 'Batch Gradient Descent: step size=0.01, l2 penalty=0.01')
```




Maybe we can get rid of the oscillations by decreasing the step size

```
In [365... w_final, cost = bgd(x_train, y_train, 0.001, 0.01, 500)
```

```
In [366... plt.figure(figsize=(7,6))
plt.plot(np.linspace(1,501,501),cost)
plt.xlabel('Iteration', fontsize=15)
plt.ylabel('Cost', fontsize=15)
plt.title('Batch Gradient Descent: step size=0.001, l2 penalty=0.01', fontsize=15)
# plt.savefig('Q3.2 0001001.png',dpi=300)
```

```
Out[366]: Text(0.5, 1.0, 'Batch Gradient Descent: step size=0.001, l2 penalty=0.01')
```


Batch Gradient Descent: step size=0.001, l2 penalty=0.01



Q3.3

In [22]: `import random`

In [215...]: `#update function for stochastic gradient descent with l2 penalty`
`def sgd_update(data, labels, weight, step, penalty):`
 `#creat the w' vector for the update and calculating the cost`
 `w_p = np.append(weight[:-1],0)`
 `return weight+step*((labels -logistic_fn(data,weight))*data - penalty*w_p)`

In [242...]: `#calculates the cost of the logistic regression cost function with l2 penalty`
`def sgd_cost(data, labels, weights, penalty):`
 `s_vals = logistic_fn(data,weights)`
 `log_s = [log(x+1e-10) for x in s_vals]`
 `log_s_comp = [log((1-x)+1e-10) for x in s_vals]`
 `return -np.dot(labels,log_s)-np.dot((1-labels),log_s_comp)+penalty*np.linalg.norm(weights)`

```
In [300... # Performs stochastic gradient descent
def sgd(data, labels, step, penalty, epochs, batch_size):

    data = add_fic(data)

    # Initialize the weight with the zero vector
    weight = np.zeros(data.shape[1])

    # Initialize the cost with the first weight vector of all zeros
    cost = [sgd_cost(data, labels, weight, penalty)]

    random.seed(42)

    for epoch in range(epochs):

        rstate = random.randint(0,data.shape[0])

        # Shuffle the data and labels for the epoch
        x_train, x_val, y_train, y_val = train_test_split(data,labels, test_size = 0.7)

        for i in range(0, x_train.shape[0], batch_size):

            # Select the mini-batch of data and labels
            x = x_train[i:i+batch_size]

            y = y_train[i:i+batch_size]

            # Compute the logistic function and update the weight
            weight = sgd_update(x[0], y, weight, step, penalty)

            # Compute the cost for the mini-batch and store it
            cost += [sgd_cost(data, labels, weight, penalty)]


    return weight, cost
```

```
In [303... weight, cost = sgd(x, labels, step=0.00001, penalty=0.01, epochs=1, batch_size=1)
```

```
In [304... plt.figure(figsize=(7,6))
plt.plot(np.linspace(1,len(cost),len(cost)),cost)
plt.xlabel('Iteration', fontsize=15)
plt.ylabel('Cost', fontsize=15)
plt.title('Stochastic Gradient Descent: step size=0.00001, 12 penalty=0.01', fontsize=15)
# plt.savefig('Q3.4 SGD.png',dpi=300)
```

```
Out[304]: Text(0.5, 1.0, 'Stochastic Gradient Descent: step size=0.00001, 12 penalty=0.01')
```

Stochastic Gradient Descent: step size=0.00001, l2 penalty=0.01



Q3.5

In [322...

```
# Performs stochastic gradient descent
def sgd_step(data, labels, delta, penalty, epochs, batch_size):

    data = add_fic(data)

    # Initialize the weight with the zero vector
    weight = np.zeros(data.shape[1])

    # Initialize the cost with the first weight vector of all zeros
    cost = [sgd_cost(data, labels, weight, penalty)]

    random.seed(42)

    for epoch in range(epochs):

        rstate = random.randint(0, data.shape[0])

        # Shuffle the data and labels for the epoch
        x_train, x_val, y_train, y_val = train_test_split(data, labels, test_size = 0.7)

        for i in range(0, x_train.shape[0], batch_size):

            # Select the mini-batch of data and labels
            x = x_train[i:i+batch_size]
```

```

y = y_train[i:i+batch_size]

# Compute the logistic function and update the weight
weight = sgd_update(x[0], y, weight, delta/(i+1), penalty)

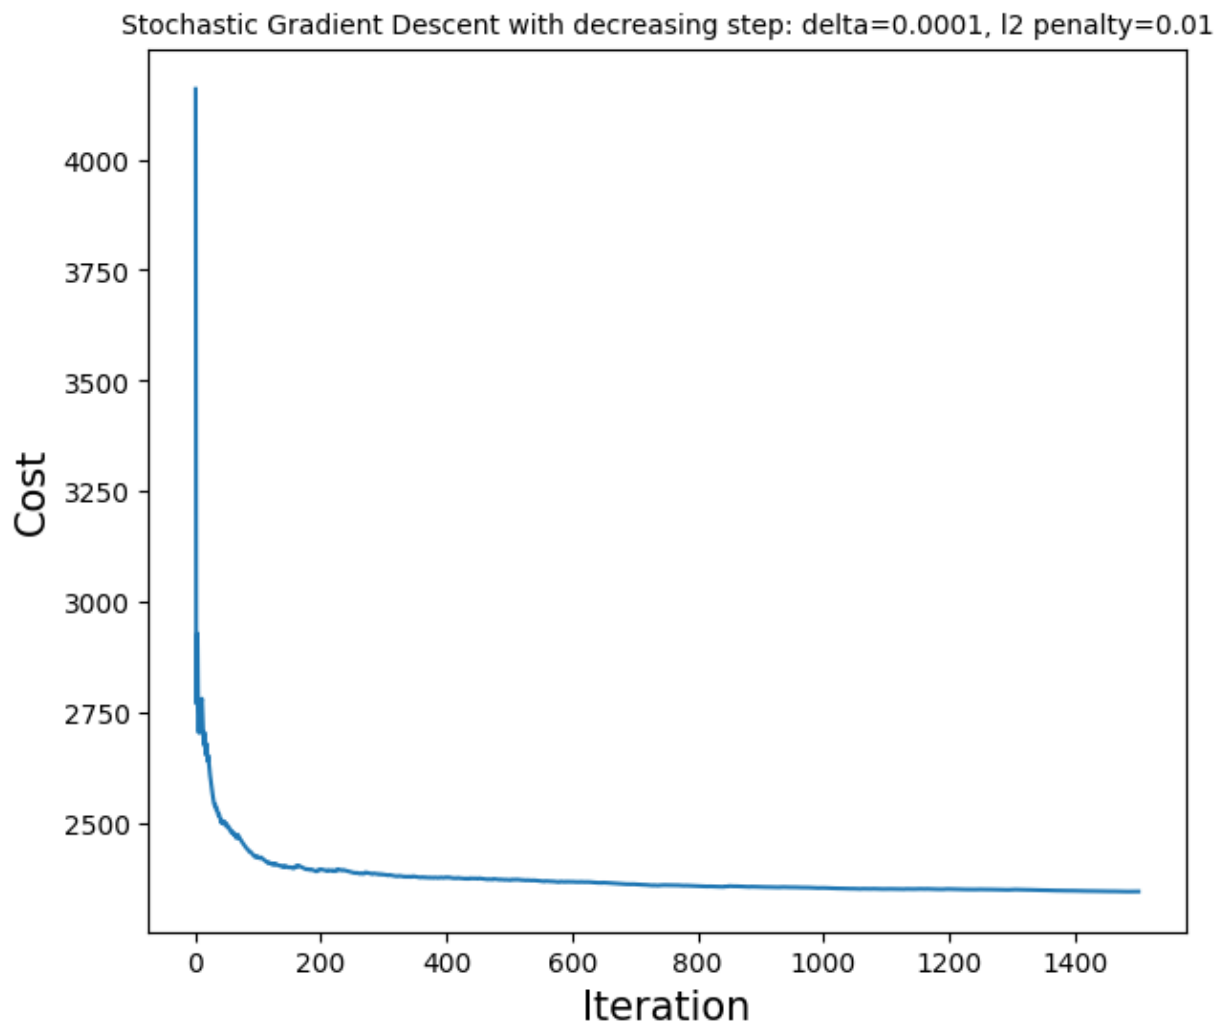
# Compute the cost for the mini-batch and store it
cost += [sgd_cost(data, labels, weight, penalty)]

return weight, cost

```

In [331... weight, cost = sgd_step(x, labels, delta=0.0001, penalty=0.01, epochs=1, batch_size=1)

In [336... plt.figure(figsize=(7,6))
plt.plot(np.linspace(1,len(cost),len(cost)),cost)
plt.xlabel('Iteration', fontsize=15)
plt.ylabel('Cost', fontsize=15)
plt.title('Stochastic Gradient Descent with decreasing step: delta=0.0001, l2 penalty=
plt.savefig('Q3.5 SGD.png',dpi=300)



Q3.6


```
In [369... #normalize the data
fit = get_fit(x)
x_train = transform(x,fit)
x_train = add_fic(x_train) #add column of 1s for the bias
```

```
In [370... w_final, cost = bgd(x_train, labels, 0.001, 0.01, 500)
```

```
In [371... plt.figure(figsize=(7,6))
plt.plot(np.linspace(1,len(cost),len(cost)),cost)
plt.xlabel('Iteration', fontsize=15)
plt.ylabel('Cost', fontsize=15)
plt.title('Batch Gradient Descent on all training: step size=0.01, l2 penalty=1', font
# plt.savefig('Q3.2 0011.png',dpi=300)
```

```
Out[371]: Text(0.5, 1.0, 'Batch Gradient Descent on all training: step size=0.01, l2 penalty=1')
```

Batch Gradient Descent on all training: step size=0.01, l2 penalty=1



```
In [375... #normalize the test data
x_test = transform(test,fit)
x_test = add_fic(x_test) #add column of 1s for the bias
```

```
In [376... y_pred = predict(x_test, w_final)
```

```
In [377... y_pred
```

```
Out[377]: array([0, 1, 1, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0,
      0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0,
      0, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0,
      0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0,
      1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
      1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1,
      0, 1, 1, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0,
      0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0,
      1, 0, 1, 0, 1, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1,
      0, 1, 0, 0, 1, 0, 1, 1, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,
      0, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 1, 1, 1, 1, 0,
      0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0,
      1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0,
      0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0,
      0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 1, 0, 0, 1, 0,
      0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
      1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 1, 1, 0,
      1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0,
      1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
      0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1,
      0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
      0, 1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
      0, 0, 0, 0, 0, 1, 1, 1, 1, 0, 0, 1, 0])
```

```
In [378... df = pd.DataFrame({'Id': np.linspace(1,len(y_pred),len(y_pred),dtype=int), 'Category':
df.to_csv('wine_preds.csv',index=False)
```

Q5.1

```
In [395... def grid_gen(xlim,ylim, num):

    # Define the range of x and y values
    x_min, x_max = -xlim, xlim
    y_min, y_max = -ylim, ylim
    num_points = num # number of points in each direction

    # Create a 1D array of x values and y values
    x = np.linspace(x_min, x_max, num_points)
    y = np.linspace(y_min, y_max, num_points)

    # Create a 2D grid of points using the meshgrid function
    X, Y = np.meshgrid(x, y)

    return X,Y
```

```
In [398... def f_eval(X,Y,p):

    return (np.abs(X)**p + np.abs(Y)**p)**(1/p)
```

```
In [408... def contour(p,xlim=None,ylim=None,num=None,f=None,grid=None,flag=True):

    if flag==True:
        X,Y = grid_gen(xlim,ylim,num)
        norm = f_eval(X,Y,p)

    else:
```

```


norm=f

fig = plt.figure(figsize=(10,10))
cs = plt.contour(X,Y,norm)
plt.clabel(cs,inline_spacing=5, fontsize=7)
plt.xlabel('X')
plt.ylabel('Y')
plt.title('Contour plot for l_{ } norm'.format(p))
plt.savefig('l_{ }_norm_plot.png'.format(p), dpi=300)

```


In [410...

```
contour(0.5,4,4,100)
```




In [411...

```
contour(1,4,4,100)
```



In [412...

```
contour(2,4,4,100)
```

```
In [413... import sys
print("Python version:", sys.version)
Python version: 3.7.15 (default, Nov 24 2022, 18:44:54) [MSC v.1916 64 bit (AMD64)]
```

```
In [414... print("NumPy version:", np.__version__)
NumPy version: 1.21.5
```

```
In [416... import sklearn
print("scikit-learn version:", sklearn.__version__)
scikit-learn version: 1.0.2
```

```
In [ ]:
```