

```
In [129... # You may want to install "gprof2dot"
import io
from collections import Counter

import numpy as np
import pandas as pd
import scipy.io
import sklearn.model_selection
import sklearn.tree
from numpy import genfromtxt
from scipy import stats
from sklearn.base import BaseEstimator, ClassifierMixin
from math import log2 as log

import pydot
```

```
In [319... eps = 1e-5 # a small number

def entropy(y):

    if len(y) == 0:
        return 0
    #assumes labels are either one or zero
    pc = sum(y)/len(y)
    pd = 1 - pc

    return -(pc*log(pc+eps)+pd*log(pd+eps))
```

```
In [321... class DecisionTree:
    def __init__(self, max_depth=3, feature_labels=None):
        self.max_depth = max_depth
        self.features = feature_labels
        self.left, self.right = None, None # for non-leaf nodes
        self.split_idx, self.thresh = None, None # for non-leaf nodes
        self.data, self.pred = None, None # for leaf nodes

    @staticmethod
    def information_gain(X, y, thresh):
        # TODO: implement information gain function

        lc = [y[i] for i in range(len(X)) if X[i]>=thresh]
        rc = [y[i] for i in range(len(X)) if X[i]<thresh]

        hs = entropy(y)
        hafter = (len(lc)*entropy(lc)+len(rc)*entropy(rc))/len(y)

        return hs-hafter

    @staticmethod
    def gini_impurity(X, y, thresh):
        # TODO: implement gini impurity function

        lc = [y[i] for i in range(len(X)) if X[i]>=thresh]
        rc = [y[i] for i in range(len(X)) if X[i]<thresh]

        l_gini = (1-(sum(lc)/len(lc))**2-(1-(sum(lc)/len(lc))))**2
        r_gini = (1-(sum(rc)/len(rc))**2-(1-(sum(rc)/len(rc))))**2
```

```

        return l_gini*len(lc)/len(y) + r_gini*len(rc)/len(y)

def split(self, X, y, idx, thresh):
    X0, idx0, X1, idx1 = self.split_test(X, idx=idx, thresh=thresh)
    y0, y1 = y[idx0], y[idx1]
    return X0, y0, X1, y1

def split_test(self, X, idx, thresh):
    idx0 = np.where(X[:, idx] < thresh)[0]
    idx1 = np.where(X[:, idx] >= thresh)[0]
    X0, X1 = X[idx0, :], X[idx1, :]
    return X0, idx0, X1, idx1

def fit(self, X, y):
    if self.max_depth > 0:
        print(self.max_depth)
        # compute entropy gain for all single-dimension splits,
        # thresholding with a linear interpolation of 10 values
        gains = []
        # The following logic prevents thresholding on exactly the minimum
        # or maximum values, which may not lead to any meaningful node
        # splits.
        thresh = np.array([
            np.linspace(np.min(X[:, i]) + eps, np.max(X[:, i]) - eps, num=10)
            for i in range(X.shape[1])
        ])
        for i in range(X.shape[1]):
            #passes the datapoints for a feature, the labels and a threshold value
            #all the gains on all the features if they were added as the next node
            gains.append([self.information_gain(X[:, i], y, t) for t in thresh[i],

        gains = np.nan_to_num(np.array(gains))

        self.split_idx, thresh_idx = np.unravel_index(np.argmax(gains), gains.shape)
        self.thresh = thresh[self.split_idx, thresh_idx]
        X0, y0, X1, y1 = self.split(X, y, idx=self.split_idx, thresh=self.thresh)
        if X0.size > 0 and X1.size > 0:
            self.left = DecisionTree(
                max_depth=self.max_depth - 1, feature_labels=self.features)
            self.left.fit(X0, y0)
            self.right = DecisionTree(
                max_depth=self.max_depth - 1, feature_labels=self.features)
            self.right.fit(X1, y1)
        else:
            self.max_depth = 0
            self.data, self.labels = X, y
            self.pred = stats.mode(y).mode[0]
    else:
        self.data, self.labels = X, y
        self.pred = stats.mode(y).mode[0]
    return self

def predict(self, X):
    if self.max_depth == 0:
        return self.pred * np.ones(X.shape[0])
    else:
        X0, idx0, X1, idx1 = self.split_test(X, idx=self.split_idx, thresh=self.thresh)
        yhat = np.zeros(X.shape[0])

```

```

        yhat[idx0] = self.left.predict(X0)
        yhat[idx1] = self.right.predict(X1)
        return yhat

    def __repr__(self):
        if self.max_depth == 0:
            return "%s (%s)" % (self.pred, self.labels.size)
        else:
            return "[%s < %s: %s | %s]" % (self.features[self.split_idx],
                                           self.thresh, self.left.__repr__(),
                                           self.right.__repr__())

```

In [211]...

```

class BaggedTrees(BaseEstimator, ClassifierMixin):
    def __init__(self, params=None, n=200):
        if params is None:
            params = {}
        self.params = params
        self.n = n
        self.decision_trees = [
            sklearn.tree.DecisionTreeClassifier(random_state=i, **self.params)
            for i in range(self.n)
        ]

    def fit(self, X, y):
        # TODO: implement function
        pass

    def predict(self, X):
        # TODO: implement function
        pass

class RandomForest(BaggedTrees):
    def __init__(self, params=None, n=200, m=1):
        if params is None:
            params = {}
        # TODO: implement function
        pass

class BoostedRandomForest(RandomForest):
    def fit(self, X, y):
        self.w = np.ones(X.shape[0]) / X.shape[0] # Weights on data
        self.a = np.zeros(self.n) # Weights on decision trees
        # TODO: implement function
        return self

    def predict(self, X):
        # TODO: implement function
        pass

```

In [4]:

```

def preprocess(data, fill_mode=True, min_freq=10, onehot_cols=[]):
    # fill_mode = False

    # Temporarily assign -1 to missing data
    data[data == ''] = '-1'

    # Hash the columns (used for handling strings)
    onehot_encoding = []

```

```

onehot_features = []
for col in onehot_cols:
    counter = Counter(data[:, col])
    for term in counter.most_common():
        if term[0] == '-1':
            continue
        if term[-1] <= min_freq:
            break
        onehot_features.append(term[0])
        onehot_encoding.append((data[:, col] == term[0]).astype(float))
    data[:, col] = '0'
onehot_encoding = np.array(onehot_encoding).T
data = np.hstack([np.array(data, dtype=float), np.array(onehot_encoding)])

# Replace missing data with the mode value. We use the mode instead of
# the mean or median because this makes more sense for categorical
# features such as gender or cabin type, which are not ordered.
if fill_mode:
    for i in range(data.shape[-1]):
        mode = stats.mode(data[((data[:, i] < -1 - eps) +
                                (data[:, i] > -1 + eps))[:, i]].mode[0])
        data[(data[:, i] > -1 - eps) * (data[:, i] < -1 + eps)[:, i] = mode

return data, onehot_features

```

```

In [6]: def evaluate(clf):
        print("Cross validation", sklearn.model_selection.cross_val_score(clf, X, y))
        if hasattr(clf, "decision_trees"):
            counter = Counter([t.tree_.feature[0] for t in clf.decision_trees])
            first_splits = [(features[term[0]], term[1]) for term in counter.most_common()]
            print("First splits", first_splits)

```

```

In [145... if __name__ == "__main__":
            dataset = "titanic"
            params = {
                "max_depth": 5,
                # "random_state": 6,
                "min_samples_leaf": 10,
            }
            N = 100

            if dataset == "titanic":
                # Load titanic data
                path_train = './dataset/titanic/titanic_training.csv'
                data = genfromtxt(path_train, delimiter=',', dtype=None, encoding=None)
                path_test = './dataset/titanic/titanic_test_data.csv'
                test_data = genfromtxt(path_test, delimiter=',', dtype=None, encoding=None)
                y = data[1:, -1] # Label = survived
                class_names = ["Died", "Survived"]
                labeled_idx = np.where(y != '')[0]

                y = np.array(y[labeled_idx])
                y = y.astype(float).astype(int)

                print("\n\nPart (b): preprocessing the titanic dataset")
                X, onehot_features = preprocess(data[1:, :-1], onehot_cols=[1, 5, 7, 8])
                X = X[labeled_idx, :]
                Z, _ = preprocess(test_data[1:, :], onehot_cols=[1, 5, 7, 8])

```

```

    assert X.shape[1] == Z.shape[1]
    features = list(data[0, :-1]) + onehot_features

elif dataset == "spam":
    features = [
        "pain", "private", "bank", "money", "drug", "spam", "prescription", "creat
        "height", "featured", "differ", "width", "other", "energy", "business", "n
        "volumes", "revision", "path", "meter", "memo", "planning", "pleased", "re
        "semicolon", "dollar", "sharp", "exclamation", "parenthesis", "square_brac
        "ampersand"
    ]
    assert len(features) == 32

    # Load spam data
    path_train = './dataset/spam/spam_data.mat'
    data = scipy.io.loadmat(path_train)
    X = data['training_data']
    y = np.squeeze(data['training_labels'])
    Z = data['test_data']
    class_names = ["Ham", "Spam"]

else:
    raise NotImplementedError("Dataset %s not handled" % dataset)

print("Features:", features)
print("Train/test size:", X.shape, Z.shape)

print("\n\nPart 0: constant classifier")
print("Accuracy", 1 - np.sum(y) / y.size)

# Basic decision tree
print("\n\nPart (a-b): simplified decision tree")
dt = DecisionTree(max_depth=3, feature_labels=features)
dt.fit(X, y)
print("Predictions", dt.predict(Z)[:100])

print("\n\nPart (c): sklearn's decision tree")
clf = sklearn.tree.DecisionTreeClassifier(random_state=0, **params)
clf.fit(X, y)
evaluate(clf)
out = io.StringIO()

# You may want to install "gprof2dot"
sklearn.tree.export_graphviz(
    clf, out_file=out, feature_names=features, class_names=class_names)
graph = pydot.graph_from_dot_data(out.getvalue())
# pydot.graph_from_dot_data(out.getvalue())[0].write_pdf("%s-tree.pdf" % dataset)

# TODO: implement and evaluate!

```

```
Features: ['pclass', 'sex', 'age', 'sibsp', 'parch', 'ticket', 'fare', 'cabin', 'emba
rked', 'male', 'female', 'S', 'C', 'Q']
Train/test size: (999, 14) (310, 14)
```

Part (a-b): simplified decision tree

Part (c): sklearn's decision tree

# Titanic preprocessing

```
titanic_train = pd.read_csv('dataset/titanic/titanic_training.csv')
```

|     | pclass | sex    | age  | sibsp | parch | ticket  | fare     | cabin | embarked | survived |
|-----|--------|--------|------|-------|-------|---------|----------|-------|----------|----------|
| 0   | 1.0    | female | 40.0 | 1.0   | 1.0   | 16966   | 134.5000 | E34   | C        | 1.0      |
| 1   | 3.0    | male   | 33.0 | 0.0   | 0.0   | 345780  | 9.5000   | NaN   | S        | 0.0      |
| 2   | 3.0    | male   | 3.0  | 4.0   | 2.0   | 347077  | 31.3875  | NaN   | S        | 1.0      |
| 3   | 2.0    | female | 50.0 | 0.0   | 1.0   | 230433  | 26.0000  | NaN   | S        | 1.0      |
| 4   | 3.0    | female | 16.0 | 1.0   | 1.0   | 2625    | 8.5167   | NaN   | C        | 1.0      |
| ... | ...    | ...    | ...  | ...   | ...   | ...     | ...      | ...   | ...      | ...      |
| 995 | 1.0    | male   | 54.0 | 0.0   | 0.0   | 17463   | 51.8625  | E46   | S        | 0.0      |
| 996 | 3.0    | female | NaN  | 3.0   | 1.0   | 4133    | 25.4667  | NaN   | S        | 0.0      |
| 997 | 3.0    | male   | 18.0 | 1.0   | 0.0   | 3101267 | 6.4958   | NaN   | S        | 0.0      |
| 998 | 2.0    | male   | 31.0 | 0.0   | 0.0   | 244270  | 13.0000  | NaN   | S        | 1.0      |
| 999 | 3.0    | female | 24.0 | 0.0   | 2.0   | PP 9549 | 16.7000  | G6    | S        | 1.0      |

```
In [332... #get the number and percentage of missing data points for each column
```

```
nulls = pd.DataFrame(columns=['feature', 'n null', 'percent null'])
```

```
count = 0
for col in titanic_train.columns:
    row = {'feature': col, 'n null':titanic_train[col].isnull().sum(),
          'percent null':titanic_train[col].isnull().sum()/len(titanic_train)}
    nulls.loc[count] = row
    count+=1
```

In [333... nulls

Out[333]:

|   | feature  | n null | percent null |
|---|----------|--------|--------------|
| 0 | pclass   | 1      | 0.001        |
| 1 | sex      | 1      | 0.001        |
| 2 | age      | 205    | 0.205        |
| 3 | sibsp    | 1      | 0.001        |
| 4 | parch    | 1      | 0.001        |
| 5 | ticket   | 1      | 0.001        |
| 6 | fare     | 2      | 0.002        |
| 7 | cabin    | 774    | 0.774        |
| 8 | embarked | 3      | 0.003        |
| 9 | survived | 1      | 0.001        |

The vast majority of the datapoint are missing values for the cabin feature so in this case, rather than impute values it make more sense to drop it as a feature. The rest of the features can be kept and the missing values imputed.

In [334... titanic\_train = titanic\_train.drop('cabin', axis=1)

Before imputing the categorical data should be converted to numerical data. We only need to do this for 'sex', 'ticket' and 'embarked'. 'sex' is easy: we can do 0 for male and 1 for female

In [335... *#create a copy for preprocessing*  
titanic\_proc = titanic\_train.copy()

In [336... count = 0  
for val in titanic\_proc['sex']:  
 if val == 'male':  
 titanic\_proc.loc[count,'sex'] = 0  
 count+=1  
 else:  
 titanic\_proc.loc[count,'sex'] = 1  
 count+=1

Next, I convert the ticket numbers to ints by converting any letters into their ASCII code

In [337... count=0

```

for string in titanic_proc['ticket']:

    if isinstance(string, float):
        titanic_proc.loc[count, 'ticket'] = int(new_string)
        count+=1

    else:
        new_string = ""

        for char in string:
            if char.isdigit():
                new_string += char
            else:
                new_string += str(ord(char))

        titanic_proc.loc[count, 'ticket'] = int(new_string)
        count += 1

```

Next, convert 'embarked' the following way: C=0, Q=1, S=2

```

In [338... count=0
for val in titanic_proc['embarked']:

    if val == 'C':
        titanic_proc.loc[count, 'embarked'] = 0
        count+=1

    elif val == 'Q':
        titanic_proc.loc[count, 'embarked'] = 1
        count+=1

    else:
        titanic_proc.loc[count, 'embarked'] = 2
        count+=1

```

```

In [339... titanic_proc

```

```

Out[339]:

```

|     | pclass | sex | age  | sibsp | parch | ticket     | fare     | embarked | survived |
|-----|--------|-----|------|-------|-------|------------|----------|----------|----------|
| 0   | 1.0    | 1   | 40.0 | 1.0   | 1.0   | 16966      | 134.5000 | 0        | 1.0      |
| 1   | 3.0    | 0   | 33.0 | 0.0   | 0.0   | 345780     | 9.5000   | 2        | 0.0      |
| 2   | 3.0    | 0   | 3.0  | 4.0   | 2.0   | 347077     | 31.3875  | 2        | 1.0      |
| 3   | 2.0    | 1   | 50.0 | 0.0   | 1.0   | 230433     | 26.0000  | 2        | 1.0      |
| 4   | 3.0    | 1   | 16.0 | 1.0   | 1.0   | 2625       | 8.5167   | 0        | 1.0      |
| ... | ...    | ... | ...  | ...   | ...   | ...        | ...      | ...      | ...      |
| 995 | 1.0    | 0   | 54.0 | 0.0   | 0.0   | 17463      | 51.8625  | 2        | 0.0      |
| 996 | 3.0    | 1   | NaN  | 3.0   | 1.0   | 4133       | 25.4667  | 2        | 0.0      |
| 997 | 3.0    | 0   | 18.0 | 1.0   | 0.0   | 3101267    | 6.4958   | 2        | 0.0      |
| 998 | 2.0    | 0   | 31.0 | 0.0   | 0.0   | 244270     | 13.0000  | 2        | 1.0      |
| 999 | 3.0    | 1   | 24.0 | 0.0   | 2.0   | 8080329549 | 16.7000  | 2        | 1.0      |

1000 rows × 9 columns



```
In [340...] from sklearn.impute import KNNImputer
```

```
In [341...] imputer = KNNImputer(n_neighbors=10)
```

```
In [342...] titanic_imputed = pd.DataFrame(imputer.fit_transform(titanic_proc), columns=titanic_pr
```

```
In [343...] titanic_labels=np.array(titanic_imputed['survived'])
```

```
In [344...] titanic_t_data = np.array(titanic_imputed.drop('survived', axis=1))
```

```
In [327...] titanic_t_data.shape
```

```
Out[327]: (1000, 8)
```

```
In [345...] classifier = DecisionTree()
```

```
In [346...] classifier.fit(titanic_t_data,titanic_labels)
```

```
3
2
1
1
2
1
1
```

```
-----
TypeError                                Traceback (most recent call last)
~\Anaconda3\lib\site-packages\IPython\core\formatters.py in __call__(self, obj)
    700         type_pprinters=self.type_pprinters,
    701         deferred_pprinters=self.deferred_pprinters)
--> 702         printer.pretty(obj)
    703         printer.flush()
    704         return stream.getvalue()

~\Anaconda3\lib\site-packages\IPython\lib\pretty.py in pretty(self, obj)
    392         if cls is not object \
    393             and callable(cls.__dict__.get('__repr__')):
--> 394             return _repr_pprint(obj, self, cycle)
    395
    396         return _default_pprint(obj, self, cycle)

~\Anaconda3\lib\site-packages\IPython\lib\pretty.py in _repr_pprint(obj, p, cycle)
    698     """A pprint that just redirects to the normal repr function."""
    699     # Find newlines and replace them with p.break_()
--> 700     output = repr(obj)
    701     lines = output.splitlines()
    702     with p.group():

~\AppData\Local\Temp\ipykernel_47288\3586919021.py in __repr__(self)
     96         return "%s (%s)" % (self.pred, self.labels.size)
     97     else:
--> 98         return "[%s < %s: %s | %s]" % (self.features[self.split_idx],
     99                                         self.thresh, self.left.__repr__(),
    100                                         self.right.__repr__())

TypeError: 'NoneType' object is not subscriptable
```

```
In [347... train, one_hots = preprocess(np.array(titanic_train)[:,-1], onehot_cols=[1,5,7])
```

```
In [351... classifier.fit(train,titanic_labels)
```

```
Out[351]: 3  
0.0 (1000)
```

```
In [ ]:
```